

CHAPTER 8

Inserting Data



he `INSERT` statement lets you enter data into tables and views in two ways: via an `INSERT` statement with a `VALUES` (and in MySQL with a `VALUES` or `SET`) clause and via an `INSERT` statement with a query. The `VALUES` clause takes a list of literal (strings, numbers, and dates represented as strings), expressions (return values from functions), or variable values.

Query values are results from `SELECT` statements that are subqueries (covered in Chapter 11). `INSERT` statements work with scalar, single-row, and multiple-row subqueries. The list of columns in the `VALUES` clause or `SELECT` clause of a query (a `SELECT` list) must map to the positional list of columns that defines the table. That list is found in the data dictionary or catalog. Alternatively to the list of columns from the data catalog, you can provide a named list of those columns. The named list *overrides* the positional (or *default*) order from the data catalog and must provide at least all mandatory columns in the table definition. Mandatory columns are those that are not null constrained.

MySQL and Oracle databases differ in how they implement the `INSERT` statement. MySQL supports multiple row inserts with a `VALUES` clause; Oracle does not support this. Oracle and MySQL both support default and override signatures as qualified in the ANSI SQL standards. Oracle also provides a multiple-table `INSERT` statement, which MySQL does not provide. This chapter covers how you enter data with an `INSERT` statement that is based on a `VALUES` clause or a subquery result statement, as well as multiple-table `INSERT` statements.

The `INSERT` statement has one significant limitation: its *default signature*. The default signature is the list of columns that define the table in the data catalog. The list is defined by the position and data type of columns. The `CREATE` statement defines the initial default signature, and the `ALTER` statement can change the number, data types, or ordering of columns in the default signature.

The default prototype for an `INSERT` statement allows for an optional column list that overrides the default list of columns. Like methods in object-oriented programming languages (OOPs), an `INSERT` statement without the optional column list constructs an instance (or row) of the table using the default constructor. The override constructor for a row is defined by any `INSERT` statement when you provide an optional column list. That's because it overrides the default constructor.

The generic prototype for an `INSERT` statement is confusing when it tries to capture both the `VALUES` clause and the result set from a query. Therefore, I've opted to provide two generic prototypes. The first uses the `VALUES` clause:

```
INSERT
INTO [{LOW_PRIORITY | DELAYED | HIGH_PRIORITY}] [IGNORE] table_name
[( column1, column2, column3, ...)]
VALUES
( value1, value2, value3, ...)
[ON DUPLICATE KEY
UPDATE column_name1 = expression1
[,      column_name2 = expression2
[, ...]]];
```

MySQL also supports multiple rows in the `VALUES` clause, which supports this modification to the `VALUES` clause:

```
VALUES
( value1A, value2A, value3A, ...)
, ( value1B, value2B, value3B, ...);
```

Overriding vs. Overloading

OOPs have special vocabularies. You define a class in OOPs by writing a program that outlines the rules for how to create an instance of a class and the methods available for any instance of that class.

Objects have special methods known as constructors that let you create instances. The default constructor generally has no formal parameters. When writing the code for a class, you can define constructors that override the default constructor's behavior. These user-defined constructors are known as *overriding* constructors. Objects also have methods that perform actions against the instance, and some of these methods are overloaded, which means a given method name supports different lists of formal parameters. This is known as *method overloading*.

The list of parameters in an overriding constructor is also known as the *overriding signature*, defined as something that serves to set apart or identify. The same logic makes the formal parameter list of any method a signature of that method.

The class is also known as an object *type*, which is a data type. A table in a database is an object type, because it contains a definition of what it can include. Every row is an instance of that object type.

As qualified in Chapter 1, SQL is a set-based declarative language. Declarative languages hide the implementation details while providing a means for the developer to state what should happen. The default signature to enter a row of data is read from the data catalog and compared against the list of values in a `VALUES` clause or query. The `INSERT` statement's optional column list lets you override the default signature to enter a row of data, but the database checks to ensure that you conform to any not null column-level constraints.

MySQL supports another less frequently used syntax that resembles the `UPDATE` statement. It's a named assignment prototype with a `SET` clause:

```
INSERT
  INTO [{LOW_PRIORITY | DELAYED | HIGH_PRIORITY}] [IGNORE] table_name
  SET column_name1 = expression1
  [, column_name2 = expression2
  [, ...]]
  [ON DUPLICATE KEY
    UPDATE column_name1 = expression2
    [, column_name2 = expression2
    [, ...]]];
```

The prototype of an `INSERT` statement based on a result set from a query is shown next:

```
INSERT
  INTO [{LOW_PRIORITY | DELAYED | HIGH_PRIORITY}] [IGNORE] table_name
  [( column1, column2, column3, ...)]
  ( SELECT value1, value2, value3, ...
    FROM   some_table
    WHERE  some_column = some_value )
  [ON DUPLICATE KEY
```

```
UPDATE column_name1 = expression1  
[,      column_name2 = expression2  
[, ...]]];
```

Notice that the prototype for an `INSERT` statement with the result set from a query doesn't use the `VALUES` clause at all. A parsing error occurs when the `VALUES` clause and query both occur in an `INSERT` statement.

Default signatures present a risk of data corruption through insertion anomalies, which occur when you enter bad data in tables. Mistakes transposing or misplacing values can occur more frequently with a default signature, because the underlying table structure can change. As a best practice, always use named notation by providing the optional list of values; this should help you avoid putting the right data in the wrong place.

**TIP**

Inserts should always rely on named notation to help you avoid adding data in the wrong columns.

This chapter provides examples that use the default and override syntax for `INSERT` statements in both Oracle and MySQL. The Oracle sections also support multiple-table `INSERT` statements and a `RETURNING INTO` clause, which is an extension of the ANSI SQL standard. Oracle uses the `RETURNING INTO` clause to manage large objects and some of the features of Oracle's dynamic SQL. Note that Oracle also supports a bulk `INSERT` statement, which is covered in Chapter 13 because it requires knowledge of PL/SQL.

Insert by Values

Inserting by the `VALUES` clause is the most common type of `INSERT` statement. It's most useful when interacting with single row inserts. You typically use this type of `INSERT` statement when working with data entered through end user web forms. In some cases, users can enter more than one row of data using a form, which occurs, for example, when a user places a meal order in a restaurant and the meal and drink are treated as order items. The restaurant order entry system would enter a single row in the `order` table and two rows in the `order_item` table (one for the meal and the other for the drink). How the code enters the multiple order items can differ between an Oracle and MySQL database, because MySQL supports multiple items in the `VALUES` clause, while Oracle doesn't. For example, an Oracle implementation might open a loop and process two dynamic `INSERT` statements (see Chapter 13), while a MySQL implementation might create, prepare, execute, and de-allocate a prepared statement that inserts two rows in a single `INSERT` statement (see Chapter 14).

The next two subsections explain how to use `INSERT` statements with the `VALUES` clause in an Oracle or MySQL database.

Oracle Insert by Values

Oracle supports only a single row insert through the `VALUES` clause. Multiple row inserts require an `INSERT` statement from a query. This section covers single row `INSERT` statements and the "Oracle Insert by Query" section covers multiple row inserts.

The `VALUES` clause of an `INSERT` statement accepts scalar values, such as strings, numbers, and dates. It also accepts calls to arrays, lists, or user-defined object types, which are called

flattened objects. Oracle supports `VARRAY` as arrays and nested tables as lists. They can both contain elements of a scalar data type or user-defined object type. Chapter 6 shows you how to create these data types.

The following sections discuss how you use the `VALUES` clause with scalar data types, how you convert various data types, and how you use the `VALUES` clause with nested tables and user-defined object data types.

Inserting Scalar Data Types

The basic syntax for an `INSERT` statement with a `VALUES` clause can include an optional *override signature* between the table name and `VALUES` keyword. With an override signature, you designate the column names and the order of entry for the `VALUES` clause elements. Without an override signature, the `INSERT` signature checks the definition of the table in the database catalog. The positional order of the column in the data catalog defines the positional, or default, signature for the `INSERT` statement. You can discover the structure of a table in Oracle or MySQL with the `DESCRIBE` command issued at the SQL*Plus or MySQL Monitor command line:

```
DESCRIBE table_name;
```

The semicolon is unnecessary in Oracle but required by the MySQL Monitor. Semicolons are execution commands and are covered in Chapter 2.

You'll see the following after describing the `rental` table in SQL*Plus:

Name	Null?	Type
RENTAL_ID	NOT NULL	NUMBER
CUSTOMER_ID	NOT NULL	NUMBER
CHECK_OUT_DATE	NOT NULL	DATE
RETURN_DATE		DATE
CREATED_BY	NOT NULL	NUMBER
CREATION_DATE	NOT NULL	DATE
LAST_UPDATED_BY	NOT NULL	NUMBER
LAST_UPDATE_DATE	NOT NULL	DATE

The `rental_id` column is a *surrogate key*, or an artificial numbering sequence. The combination of the `customer_id` and `check_out_date` serves as a *natural key*, because a `DATE` data type is a date-time value. If it were only a date, the customer would be limited to a single entry for each day, and limiting customer rentals to one per day isn't a good business model.

The basic `INSERT` statement would require that you look up the next sequence value before using it. You should also look up the surrogate key column value that maps to the row where your unique customer is stored in the contact table. For this example, assume the following facts:

- Next sequence value is 1086
- Customer's surrogate key value is 1009
- Current date-time is represented by the value from the `SYSDATE` function
- Return date is the fifth date from today
- User adding and updating the row has a primary (surrogate) key value of 1
- Creation and last update date are the value returned from the `SYSDATE` function.

222 Oracle Database 11g & MySQL 5.6 Developer Handbook

An INSERT statement must include a list of values that match the positional data types of the database catalog, or it must use an override signature for all mandatory columns.

You can now write the following INSERT statement, which relies on the default signature:

```
SQL> INSERT INTO rental
2  VALUES
3  ( 1086
4  , 1009
5  , SYSDATE
6  , TRUNC(SYSDATE + 5)
7  , 1
8  , SYSDATE
9  , 1
10 , SYSDATE);
```

If you weren't using SYSDATE for the date-time value on line 5, you could manually enter a date-time with the following Oracle proprietary syntax:

```
5 , TO_DATE('15-APR-2011 12:53:01', 'DD-MON-YYYY HH24:MI:SS')
```

The TO_DATE function is an Oracle-specific function. The generic conversion function would be the CAST function. The problem with a CAST function by itself is that it can't handle a format mask other than the database defaults ('DD-MON-RR' or 'DD-MON-YYYY')—for example, consider this syntax:

```
5 , CAST('15-APR-2011 12:53:02' AS DATE)
```

It raises the following error:

```
5 , CAST('15-APR-2011 12:53:02' AS DATE) FROM dual
*
ERROR at line 1:
ORA-01830: date format picture ends before converting entire input string
```

You actually need to double cast this type of format mask when you want to store it as a DATE data type. The working syntax casts the date-time string as a TIMESTAMP data type before recasting the TIMESTAMP to a DATE, like so:

```
5 , CAST(CAST('15-APR-2011 12:53:02' AS TIMESTAMP) AS DATE)
```

Before you could have written the preceding INSERT statement, you would need to run some queries to find the values. You would secure the next value from a rental_s1 sequence in an Oracle database with the following command:

```
SQL> SELECT rental_s1.NEXTVAL FROM dual;
```

This assumes two things, because sequences are separate objects from tables. First, code from which the values in a table's surrogate key column come must appear in the correct sequence. Second, a sequence value is inserted only once into a table as a primary key value.

In place of a query that finds the next sequence value, you would simply use a call against the .NEXTVAL pseudo column in the VALUES clause. You would replace line 3 with this:

```
3 ( rental_s1.NEXTVAL
```

The `.NEXTVAL` is a pseudo column, and it instantiates an instance of a sequence in the current session. After a call to a sequence with the `.NEXTVAL` pseudo column, you can also call back the prior sequence value with the `.CURRVAL` pseudo column.



NOTE

Sequences are separate objects from tables, and your code ensures that only the appropriate sequence maps to the correct table.

Assuming the following query would return a single row, you can use the `contact_id` value as the `customer_id` value in the `rental` table:

```
SQL> SELECT    contact_id
      2 FROM      contact
      3 WHERE     last_name = 'Potter'
      4 AND      first_name = 'Harry';
```

Taking three steps like this is unnecessary, however, because you can call the next sequence value and find the valid `customer_id` value inside the `VALUES` clause of the `INSERT` statement. The following `INSERT` statement uses an override signature and calls for the next sequence value on line 11. It also uses a scalar subquery to look up the correct `customer_id` value with a scalar subquery on lines 12 through 15.

```
SQL> INSERT INTO rental
      2 ( rental_id
      3 , customer_id
      4 , check_out_date
      5 , return_date
      6 , created_by
      7 , creation_date
      8 , last_updated_by
      9 , last_update_date )
     10 VALUES
     11 ( rental_s1.nextval
     12 , (SELECT    contact_id
     13     FROM      contact
     14     WHERE     last_name = 'Potter'
     15     AND      first_name = 'Harry')
     16 , SYSDATE
     17 , TRUNC(SYSDATE + 5)
     18 , 1
     19 , SYSDATE
     20 , 3
     21 , SYSDATE);
```

When a subquery returns two or more rows because the conditions in the `WHERE` clause failed to find and return a unique row, the `INSERT` statement would fail with the following message:

```
, (SELECT    contact_id
      *
ERROR at line 3:
ORA-01427: single-row subquery returns more than one row
```

In fact, the statement could fail when two or more “Harry Potter” names exist in the data set, because three columns make up the natural key of the contact table. The third column is the `member_id`, and all three should be qualified inside a scalar subquery to guarantee that it returns only one row of data.

Handling Oracle’s Large Objects Oracle’s large objects present a small problem when they’re not null constrained in the table definition. In this case, you must insert empty object containers or references when you perform an `INSERT` statement.

Assume, for example, that you have the following three large object columns in a table:

Name	Null?	Type
-----	-----	-----
ITEM_DESC	NOT NULL	CLOB
ITEM_ICON	NOT NULL	BLOB
ITEM_PHOTO		BINARY FILE LOB

The `item_desc` column uses a CLOB (Character Large Object) data type, and it is a required column; it could hold a lengthy description of a movie, for example. The `item_icon` column uses a BLOB (Binary Large Object) data type, and it is also a required column. It could hold a graphic image. The `item_photo` column uses a binary file (an externally managed file). It could hold a null or a reference to an external graphic image; fortunately, the `item_photo` column isn’t a required column.

Oracle provides two functions that let you enter an empty large object:

```
empty_blob()
empty_clob()
```

Although you could insert a null value in the `item_photo` column, you can also enter a reference to an Oracle database virtual directory file. Here’s the syntax to enter a valid `BFILE` name with the `BFILENAME` function call:

```
10  , BFILENAME('VIRTUAL_DIRECTORY_NAME', 'file_name.png')
```

You can insert a large character or binary stream into BLOB and CLOB data types by using the stored procedures and functions available in the `DBMS_LOB` package. These operations require a working knowledge of PL/SQL programming, which isn’t covered until Chapter 13, where you’ll see an example that inserts a large CLOB column.

You can use an `EMPTY_CLOB` function or a string literal up to 32,767 bytes long in a `VALUES` clause. You must use the `DBMS_LOB` package when you insert a string that is longer than 32,767 bytes. That also changes the nature of the `INSERT` statement and requires that you append the `RETURNING INTO` clause. Here’s the prototype for this Oracle proprietary syntax:

```
INSERT INTO some_table
[( column1, column2, column3, ...)]
VALUES
( value1, value2, value3, ...)
RETURNING column1 INTO local_variable;
```

The `local_variable` is a reference to a procedural programming language. It lets you insert a character stream into a target CLOB column or a binary stream into a BLOB column.

DBA Heads-up on Large Object Storage

CLOB and BLOB columns are stored with the rest of a row when they’re smaller than 4000 bytes. Larger versions are stored out-of-line, which means they’re placed in a contiguous space that is away from the rest of the related row.

DBAs often designate a special tablespace for the storage clauses of BLOB and CLOB columns. This extra step is beneficial, because large objects change less frequently and consume a lot of storage. They’re also generally on different backup schedules than other transactional columns in a table.

Capturing the Last Sequence Value Sometimes you insert into a series of tables in the scope of a transaction, like those described in Chapter 4. In this scenario, one table gets the new sequence value (with a call to `sequence_name.NEXTVAL`) and enters it as the surrogate primary key, and another table needs a copy of that primary key to enter into a foreign key column. While scalar subqueries can solve this problem, Oracle provides the `.CURRVAL` pseudo column for this purpose.

The steps to demonstrate this behavior require a parent and child table. The parent table is defined as follows:

Name	Null?	Type
-----	-----	-----
PARENT_ID	NOT NULL	NUMBER
PARENT_NAME		VARCHAR2 (10)

The `parent_id` column is the primary key for the parent table. You include the `parent_id` column in the child table. In the child table, the `parent_id` column holds a copy of a valid primary key column value as a foreign key to the parent table.

Name	Null?	Type
-----	-----	-----
CHILD_ID	NOT NULL	NUMBER
PARENT_ID		NUMBER
PARENT_NAME		VARCHAR2 (10)

After creating the two tables, you can manage inserts into them with the `.NEXTVAL` and `.CURRVAL` pseudo columns. The sequence calls with the `.NEXTVAL` insert primary key values, and the sequence calls with the `.CURRVAL` insert foreign key values.

You would perform these two `INSERT` statements as a group:

```
SQL> INSERT INTO parent
2  VALUES
3  ( parent_s1.NEXTVAL
4  , 'One Parent');

SQL> INSERT INTO child
2  VALUES
3  ( child_s1.NEXTVAL
4  , parent_s1.CURRVAL
5  , 'One Child');
```

226 Oracle Database 11g & MySQL 5.6 Developer Handbook

The `.CURRVAL` pseudo column for any sequence fetches the value placed in memory by call to the `.NEXTVAL` pseudo column. Any attempt to call the `.CURRVAL` pseudo column before the `.NEXTVAL` pseudo column raises an ORA-02289 exception. The text message for that error says the sequence doesn't exist, which actually means that it doesn't exist in the scope of the current session. Line 4 in the insert into the `child` table depends on line 3 in the insert into the `parent` table.

You can use comments in `INSERT` statements to map to columns in the table. For example, the following shows the technique for the `child` table from the preceding example:

```
SQL> INSERT INTO child
2  VALUES
3  ( child_s1.NEXTVAL      -- CHILD_ID
4  , parent_s1.CURRVAL    -- PARENT_ID
5  , 'One Child')         -- CHILD_NAME
6  /
```

Comments on the lines of the `VALUES` clause identify the columns where the values are inserted. A semicolon doesn't execute this statement, because a trailing comment would trigger a runtime exception. You must use the semicolon or forward slash on the line below the last `VALUES` element to include the last comment.

TIP

A comment on the last line of any statement requires that you exclude the semicolon and place it or a forward slash on the next line.

Data Type Conversions

Oracle supports a series of conversion functions that let you convert data types from one type to another. The generic SQL conversion function is `CAST`, which lets you convert the following data types.

Convert from `BINARY_FLOAT` or `BINARY_DOUBLE` Data Type to `BINARY_FLOAT`, `BINARY_DOUBLE`, `CHAR`, `VARCHAR2`, `NUMBER`, `DATE`, `TIMESTAMP`, `NCHAR`, `NVARCHAR`

Convert from `CHAR` or `VARCHAR2` Data Type to `BINARY_FLOAT`, `BINARY_DOUBLE`, `CHAR`, `VARCHAR2`, `NUMBER`, `DATE`, `TIMESTAMP`, `DATE`, `TIMESTAMP`, `INTERVAL`, `RAW`, `ROWID`, `UROWID`, `NCHAR`, `NVARCHAR`

Here's an example of converting a string literal date into a timestamp:

```
CAST('14-FEB-2011' AS TIMESTAMP WITH LOCAL TIME ZONE)
```

This example works because the date literal conforms to the default format mask for a date in an Oracle database. A nonconforming date literal would raise a conversion error. Many possibilities are available because you can organize the valid elements of dates many ways. A nonconforming date literal should be converted by using the `TO_DATE` or `TO_TIMESTAMP` function, because each of these lets you specify an overriding date format mask value, such as this conversion to a `DATE` data type:

```
TO_DATE('2011-02-14', 'YYYY-MM-DD')
```

Or this conversion to a `TIMESTAMP` data type:

```
TO_TIMESTAMP('2011-02-14 18:11:28.1500', 'YYYY-MM-DD HH24:MI:SS.FF')
```

Converting to an `INTERVAL` data type is covered in the next section, because you first must extract a time property as a number. It's also possible that implicit casting of a numeric string can change the base data type to an integer for you. The method of implicit or explicit conversion depends on how you get the initial data value.

Convert from `NUMBER` Data Type to `BINARY_FLOAT`, `BINARY_DOUBLE`, `CHAR`, `VARCHAR2`, `NUMBER`, `DATE`, `TIMESTAMP`, `NCHAR`, `NVARCHAR`

Interval conversions are a bit more complex, because you need more than one function to convert them. Typically, you pull the value from a `DATE` or `TIMESTAMP` data type and `EXTRACT` the element of time by identifying its type before converting that value to an `INTERVAL` type. The following provides an example:

```
NUMTODSINTERVAL(EXTRACT(MINUTE FROM some_date), 'MINUTE')
```

You will use this type of built-in function layering frequently in some situations. It's always a better approach to understand and use the built-in functions before you write your own stored functions.

Convert from `DATETIME` or `INTERVAL` Data Type to `CHAR`, `VARCHAR2`, `DATE`, `TIMESTAMP`, `DATE`, `TIMESTAMP`, `INTERVAL`, `NCHAR`, `NVARCHAR`

Convert from `RAW` Data Type to `CHAR`, `VARCHAR2`, `RAW`, `NCHAR`, `NVARCHAR`

Convert from `ROWID` or `UROWID` Data Type to `CHAR`, `VARCHAR2`, `ROWID`, `UROWID`, `NCHAR`, `NVARCHAR`

NOTE

You cannot cast a `UROWID` to a `ROWID` in the `UROWID` of an index organized table.

Convert from `NCHAR` or `NVARCHAR2` Data Type to `BINARY_FLOAT`, `BINARY_DOUBLE`, `NUMBER`, `NCHAR`, `NVARCHAR`

Inserting Arrays and Nested Tables

The ability to insert arrays and nested tables in an Oracle database is an important feature made possible by the *object-relational* technology of the database. You can access these embedded structures only through the containing table, which makes them much like inner classes in object-oriented programming. From a database modeling perspective, they're ID-dependent data sets, because the only relationship is through the row of the containing table.

You can walk through a simple design and development exercise by creating a collection of a scalar data type, a table that contains the data type, and then an `INSERT` statement to populate the table with data. You create the user-defined collection data type by using this syntax:

```
CREATE TYPE number_array IS VARRAY(10) OF NUMBER;
```

228 Oracle Database 11g & MySQL 5.6 Developer Handbook

After you have the user-defined collection type, create a table that uses it and a sequence for an automatic numbering column, like so:

```
SQL> CREATE TABLE sample_nester
  2  ( nester_id      NUMBER
  3    , array_column NUMBER_ARRAY);
SQL> CREATE SEQUENCE sample_nester_s1;
```

You enter values into the table by calling a collection constructor. Calls are made to the data type name, not the column name, like so:

```
SQL> INSERT INTO sample_nester
  2  VALUES
  3  ( sample_nester_s1.NEXTVAL
  4    , NUMBER_ARRAY(0,1,2,3,4,5,6,7,8,9));
```

Here are formatting instructions and an ordinary query against this table:

```
SQL> COLUMN ARRAY_COLUMN FORMAT A44
SQL> SELECT * from sample_nester;
```

which returns the following:

```
NESTER_ID  ARRAY_COLUMN
-----
1 NUMBER_ARRAY(0, 1, 2, 3, 4, 5, 6, 7, 8, 9)
```

The value in the `ARRAY_COLUMN` is a call to the user-defined collection data type's constructor. This collection is a simple array of ten numbers. You can join the `nester_id` column against all ten elements in the collection with the following syntax:

```
SQL> SELECT nester_id
  2    ,      collection.column_value
  3  FROM    sample_nester CROSS JOIN TABLE(array_column) collection;
```

The `TABLE` function extracts the element of the collection into a SQL result set, which can then be treated like a normal set of rows from any table. Oracle always returns collections of base scalar data types into a `COLUMN_VALUE` column. The results from the sample query follow:

```
NESTER_ID  COLUMN_VALUE
-----
1          0
1          1
1          2
...
1          8
1          9
```

Here's a multilevel `INSERT` statement based on the employee table from Chapter 6. The following inserts a single row that contains an `address_list` collection of two instances of the

`address_type` user-defined object type, which in turn holds a collection of a `street_list` nested table of variable-length strings:

```
SQL> INSERT INTO employee
  2  ( employee_id
  3    , first_name
  4    , last_name
  5    , home_address )
  6  VALUES
  7  ( employee_s1.NEXTVAL
  8    , 'Sam'
  9    , 'Yosemite'
 10    , address_list(
 11      address_type( 1
 12                    , street_list('1111 Broadway','Suite 322')
 13                    , 'Oakland'
 14                    , 'CA'
 15                    , '94612')
 16    , address_type( 2
 17                    , street_list('1111 Broadway','Suite 525')
 18                    , 'Oakland'
 19                    , 'CA'
 20                    , '94612')));
```

Lines 10 through 20 insert a nested table (list) of two `address_type` instances. The sequence numbers are manually entered because this type of design would always start elements of a nested table with a sequence value of 1. This is an implementation of an ID-dependent relationship. The nested table is accessible only through the row of a table, and as such, it acts only when connected with the containing row.

NOTE

Nested tables are complex to access and model. As mentioned in Chapter 6, they also create type chaining, which presents maintenance headaches during major software releases. If you opt to use them, you should have a good reason for adding the complexity.

Multiple-Table Insert Statements

As mentioned, Oracle SQL syntax lets you perform multiple-table inserts with the `INSERT` statement. A multiple-table `INSERT` statement can be useful when you receive an import source file that belongs in more than one table. You should be aware of some caveats with this type of multiple-table insert. For example, a one-to-one mapping must exist between all the data in the same row. When you have a one-to-many relationship between columns in the single import source table, the `MERGE` statement is a better solution, as discussed in Chapter 12.

Here's the prototype for the multiple table `INSERT` statement:

```
INSERT {ALL | FIRST}
[WHEN comparison_clause THEN ]
INSERT INTO table_name_one
( column_list )
```

230 Oracle Database 11g & MySQL 5.6 Developer Handbook

```
VALUES
( value_list )
[[WHEN comparison_clause THEN ]
INSERT INTO table_name_two
( column_list )
VALUES
( value_list ) ]
[ ... ]
[ELSE
INSERT INTO table_name_else
( column_list )
VALUES
( value_list ) ]
query_statement;
```

A multiple-table insert can be performed in three ways: One uses the `ALL` keyword but excludes the `WHEN` clauses. The second uses the `ALL` keyword while including the `WHEN` clauses. The last uses a `FIRST` keyword instead of the `ALL` keyword and inserts to the first table before moving to the second, and so forth.

The multiple-table `INSERT` statement variations are discussed in the following subsections, but they're all supported by the following three tables and sequences.

The `rank_index` is the first table, and it contains a string for the military service, such as Army, Navy, Marines, and Air Force, with their accompanying abbreviated and full-titled ranks.

```
-- Create the rank_index table and sequence.
CREATE TABLE rank_index
( rank_id          NUMBER
, rank_service     VARCHAR2(10)
, rank_short_name  VARCHAR2(4)
, rank_full_name   VARCHAR2(30));
CREATE SEQUENCE rank_index_s;
```

The `soldier` and `sailor` tables are target tables for the inserts from the multiple-table `INSERT` statement. They contain the rank and name of soldiers and sailors, respectively:

```
-- Create the soldier table and sequence.
CREATE TABLE soldier
( soldier_id       NUMBER
, soldier_rank     VARCHAR2(4)
, soldier_name     VARCHAR2(20));
CREATE SEQUENCE soldier_s;

-- Create the sailor table and sequence.
CREATE TABLE sailor
( sailor_id        NUMBER
, sailor_rank      VARCHAR2(4)
, sailor_name      VARCHAR2(20));
CREATE SEQUENCE sailor_s;
```

All the examples get their data from a `SELECT` statement that uses the dual pseudo table and a fabricated result set. The `INSERT ALL` or `INSERT FIRST` statement inserts the data into one

or both of the target tables. The tables are unconstrained because constraints aren't required for the examples.

Multiple-Table INSERT ALL Without WHEN Clauses Our first example shows you how to insert data into multiple tables without any qualifications. Values from each row returned by the query are inserted into the `soldier` and `sailor` tables. Unfiltered multiple-table `INSERT` statements put rows into all tables referenced by an `INTO` clause. In this regard, the `INSERT ALL` statement works like a *switch* statement with fall through in C#, C++, or Java. (*Fall through* means that after meeting the condition of one case statement, all subsequent case statements are valid and their code blocks also run.)

```
SQL> INSERT ALL
  2 INTO soldier
  3 VALUES
  4 (soldier_s.NEXTVAL,service_rank,service_member_name)
  5 INTO sailor VALUES
  6 (sailor_s.NEXTVAL,service_rank,service_member_name)
  7 SELECT 'MSG' AS service_rank
  8 ,      'Ernest G. Bilko' AS service_member_name FROM dual
  9 UNION ALL
 10 SELECT 'CPO' AS service_rank
 11 ,      'David Vaught' AS service_member_name FROM dual;
```

No `ELSE` block is used in the example. This type of statement would also run the `ELSE` block and perform any insert found in it.

Queries against the target tables show you that both rows are inserted into both tables:

```
SQL> SELECT * FROM soldier;

SOLDIER_ID SOLD SOLDIER_NAME
-----
1 MSG Ernest G. Bilko
2 CPO David Vaught

SQL> SELECT * FROM sailor;

SAILOR_ID SAIL SAILOR_NAME
-----
1 MSG Ernest G. Bilko
2 CPO David Vaught
```

This type of statement is useful when you want to put data from one row of a table or view into multiple tables. Unfiltered `INSERT ALL` statements don't let you choose among a set of tables (like a filtered `INSERT ALL` statement), and they're used less than filtered statements.

Multiple-Table INSERT ALL with WHEN Clauses The multiple-table `INSERT ALL` statement also works with `WHEN` clauses that determine into which table they'll insert. The logic can include subqueries, as shown in the following example:

```
SQL> INSERT ALL
  2 WHEN service_rank IN (SELECT rank_short_name
  3                        FROM rank_index
  4                        WHERE rank_service = 'ARMY') THEN
```

232 Oracle Database 11g & MySQL 5.6 Developer Handbook

```

5  INTO soldier
6  VALUES
7  (soldier_s.NEXTVAL,service_rank,service_member_name)
8  WHEN service_rank IN (SELECT rank_short_name
9                        FROM   rank_index
10                       WHERE  rank_service = 'NAVY') THEN
11 INTO sailor
12 VALUES
13 (sailor_s.NEXTVAL,service_rank,service_member_name)
14 SELECT 'MSG' AS service_rank
15 ,      'Ernest G. Bilko' AS service_member_name FROM dual
16 UNION ALL
17 SELECT 'CPO' AS service_rank
18 ,      'David Vaught' AS service_member_name FROM dual;

```

The WHEN clause on line 2 checks whether the service rank belongs in the Army. It inserts into the `soldier` table any row in which the query's `service_rank` value matches the subquery's `rank_short_name` value. The second WHEN clause does the same kind of evaluation against Navy ranks. Any rows that don't match one of the two criteria are discarded because there is no ELSE clause.

Queries against the target tables yield the following results:

```

SQL> SELECT * FROM soldier;

SOLDIER_ID SOLD  SOLDIER_NAME
-----
1 MSG      Ernest G. Bilko

SQL> SELECT * FROM sailor;

SAILOR_ID SAIL  SAILOR_NAME
-----
2 CPO      David Vaught

```

The filtered `INSERT ALL` places rows from one source query into one or only the correct tables. This is the best practice—or at least the most frequently used version of the statement.

Multiple-Table INSERT FIRST with WHEN Clauses The `INSERT FIRST` works differently from the `INSERT ALL` statement. The `INSERT FIRST` inserts data into the first table only when it meets a WHEN clause condition. This means it performs like a switch statement in C#, C++, or Java where fall through is disabled. For your reference (in case you don't write programs in those languages), you disable fall through by adding a break statement in each case statement's code block. The break statement signals completion and forces an exit from the switch statement. The `FIRST` keyword effectively does that for all WHEN clause statement blocks.

Here's an example using the concept of conscripts (draftees). The first one goes to the Army (line 2), the next four go to the Navy (line 6), and any others get to go home without serving in the military:

```

SQL> INSERT FIRST
2  WHEN id < 2 THEN
3  INTO soldier

```



```

4  VALUES
5  (soldier_s.NEXTVAL, 'PVT', draftee)
6  WHEN id BETWEEN 2 AND 5 THEN
7  INTO sailor
8  VALUES
9  (sailor_s.NEXTVAL, 'SR', draftee)
10 SELECT 1 AS ID, 'John Sanchez' AS draftee FROM dual
11 UNION ALL
12 SELECT 2 AS ID, 'Michael Deegan' AS draftee
13 FROM dual
14 UNION ALL
15 SELECT 3 AS ID, 'Jon Voight' AS draftee FROM dual;

```

You'll see in the result set that only two draftees went in the Navy, so there weren't enough conscripts drafted today. The statement will need to be rewritten tomorrow for the new batch of draftees unless the rules change every day. The queries and results are shown next:

```
SQL> SELECT * FROM soldier;
```

```

SOLDIER_ID SOLD SOLDIER_NAME
-----
1 PVT   John Sanchez

```

```
SQL> SELECT * FROM sailor;
```

```

SAILOR_ID SAIL SAILOR_NAME
-----
2 SR     Michael Deegan
3 SR     Jon Voight

```

My suggestion is that `INSERT FIRST` statement is probably suited to dynamic creation inside Native Dynamic SQL (NDS). You can see examples of NDS in Chapter 13.

MySQL Insert by Values

`INSERT` statements in the MySQL database are used similarly to how they're used in the Oracle database. MySQL does let you insert multiple rows with a single `VALUES` clause, which isn't supported by the Oracle database. On the other hand, MySQL is a relational database and doesn't support object types or nested tables, as does Oracle 11g.

The following sections cover inserting scalar data types and performing data type conversions. MySQL treats all data types, including large character and binary strings, as scalar data types.

Inserting Scalar Data Types

MySQL lets you insert one or more rows through a `VALUES` clause. You don't have to call a sequence to use the clause during inserts, because sequences are a table property in MySQL. The `INSERT` statement mirrors more or less that of an Oracle database, because `INSERT` statements adhere to the ANSI SQL standard.

You can include an overriding signature in a MySQL `INSERT` statement just as you can in Oracle. You can exclude the single column designed with the `AUTO_INCREMENT` phrase, as well as any optional columns. The column created with an `AUTO_INCREMENT` phrase is a *surrogate*

primary key column. Like Oracle, the `INSERT` statement checks the definition of the table in the data catalog. It checks the catalog when you exclude a column list and verifies constraints when you attempt to insert data.

Mirroring the steps from the Oracle discussion, you can use the `DESCRIBE` command to view the table definition from the catalog. The semicolon is required when describing a table in the MySQL Monitor, like so:

```
DESCRIBE table_name;
```

Using our `rental` table in My SQL requires a couple data type changes. That’s because MySQL’s `DATE` is a date, not a date-time data type. The `DATETIME` data type in MySQL is a date-time type equivalent to the `DATE` data type in an Oracle database. Here’s the definition of the `rental` table (without the default value column) in MySQL:

Field	Type	Null	Key	Extra
rental_id	int(10) unsigned	NO	PRI	auto_increment
customer_id	int(10) unsigned	NO	MUL	
check_out_date	time	NO		
return_date	datetime	NO		
created_by	int(10) unsigned	NO	MUL	
creation_date	time	NO		
last_updated_by	int(10) unsigned	NO	MUL	
last_update_date	time	NO		

Like the Oracle table, the `rental_id` column is a surrogate primary key. The extra column description tells you that it’s an automatically generated sequence value. Sequences are positive integers or doubles, which means they’re defined as unsigned integers or unsigned doubles. It’s a convention, and a recommendation, that surrogate primary keys use integers until they approach the 4 billion limit of the data type.

The syntax for this `INSERT` statement with literal values doesn’t differ much from the Oracle equivalent, except for the MySQL proprietary `NOW()`, `ADDDATE()`, and `UTC_DATE()` function calls.

```
INSERT INTO rental
VALUES
( 1
, 2
, NOW()
, ADDDATE(UTC_DATE(),INTERVAL 5 DAY)
, 1
, NOW()
, 1
, NOW());
```

You can discover the next sequence value by using this syntax in MySQL Monitor:

```
show table status like 'rental'\G
```

The \G displays output vertically with the column names on the left and values on the right, as covered in Chapter 2. The show command displays this after one insert into the table:

```
***** 1. row *****
      Name: rental
      Engine: InnoDB
      Version: 10
      Row_format: Compact
      Rows: 1
      Avg_row_length: 16384
      Data_length: 16384
      Max_data_length: 0
      Index_length: 49152
      Data_free: 8388608
      Auto_increment: 2
      Create_time: 2011-01-14 22:18:59
      Update_time: NULL
      Check_time: NULL
      Collation: latin1_swedish_ci
      Checksum: NULL
      Create_options:
      Comment:
```

The important difference between the surrogate primary key column in Oracle and the same in MySQL deals with the auto-incrementing columns in MySQL. Although you do need to provide the column placeholder in the VALUES clause in MySQL, you can simply enter NULL and the correct sequence will be entered. This isn't true in Oracle, unless you've deployed an ON INSERT trigger to capture the sequence value and mapped it to the rental_id column.

In short, here is the preferred way to write an INSERT statement that uses the default signature and literal values or expression results (such as the dates from the NOW function or the date from the ADDDATE function). The INSERT statement uses a VALUES clause in MySQL:

```
INSERT INTO rental
VALUES
( null
, 2
, NOW()
, ADDDATE(UTC_DATE(), INTERVAL 5 DAY)
, 1
, NOW()
, 1
, NOW());
```

At this point, you know how to perform an INSERT statement with a null value and the default signature. It is preferred that you use a list of columns to designate what you're inserting in the VALUES clause. An INSERT statement that uses an overriding signature that excludes the auto-incrementing column follows:

```
INSERT INTO rental
( customer_id
, check_out_date
```

236 Oracle Database 11g & MySQL 5.6 Developer Handbook

```
, return_date
, created_by
, creation_date
, last_updated_by
, last_update_date )
VALUES
( SELECT    c.contact_id
  FROM      contact c JOIN member m
  ON        c.member_id = m.member_id
  WHERE     c.first_name = 'Harry'
  AND       IFNULL(c.middle_name,'x') = IFNULL(null,'x')
  AND       c.last_name = 'Potter'
  AND       m.account_number = 'SLC-000006' )
, NOW()
, ADDDATE(UTC_DATE(),INTERVAL 5 DAY)
, 1
, NOW()
, 1
, NOW());
```



NOTE

Remember that you can use scalar subqueries in MySQL, but they do not exist in Oracle. Subqueries in the VALUES clause can return only a single row when matched with literal values.

You can also rewrite this in MySQL to work with the SET clause. The only problem with the SET clause, however, is that it's not portable to other databases. Here's an example of the syntax to INSERT with the SET clause:

```
INSERT INTO rental
SET customer_id = ( SELECT c.contact_id
                   FROM   contact c JOIN member m
                   ON     c.member_id = m.member_id
                   WHERE  c.first_name = 'Harry'
                   AND     IFNULL(c.middle_name,'x') = IFNULL(null,'x')
                   AND     c.last_name = 'Potter'
                   AND     m.account_number = 'SLC-000006' )
, check_out_date = NOW()
, return_date = ADDDATE(UTC_DATE(),INTERVAL 5 DAY)
, created_by = 1
, creation_date = NOW()
, last_updated_by = 1
, last_update_date = NOW();
```

The `customer_id` value comes from a scalar subquery (see Chapter 11) against the `contact` and `member` tables. The `WHERE` clause contains the business elements that help identify the surrogate primary key (the `contact_id` column) for the `contact` table.

Data Type Conversions

Like the Oracle database, MySQL supports the `CAST` and `CONVERT` functions and the `BINARY` operator. The `CAST` function works with the following data types: `BINARY`, `CHAR`, `DATE`, `DATETIME`, `SIGNED` or `UNSIGNED INTEGER`, and `TIME`. The `CONVERT` function lets you modify a string's character set, and the `BINARY` operator is shorthand for a `CAST` function that changes a character to a binary string.

CAST Function The `CAST` function also lets you specify a character set that you want to apply through the casting operation. Here's the prototype for the `CAST` function:

```
CAST(expression_or_variable AS data_type [CHARACTER SET character_set])
```

You would convert a double precision number to a string like so:

```
CAST(111.4586 AS CHAR(30))
```

Alternatively, you can change the character set when you cast a double to a character, like this:

```
CAST(111.4586 AS CHAR(30) CHARACTER SET utf8)
```

NOTE

Casting a string to a `DATE`, `DATETIME`, or `TIME` data type has some added restrictions. The string must conform to the default string literal for a date, which is `YYYY-MM-DD` (a four-digit year, two-digit month, and two-digit day). Any attempt to convert a different format mask fails and returns a null value.

CONVERT Function The `CONVERT` function is slightly different from the `CAST` function. Here's the prototype:

```
CONVERT([character_set] string_value USING character_set)
```

To convert a string from the default character set to a Unicode character set, use the following syntax:

```
CONVERT('Hello mate!' USING utf8)
```

BINARY Operator The difference between a function and operator in this case is parentheses: there aren't any in the operator. Here's the syntax to convert a string to a binary string:

```
BINARY 'Hello Mate!'
```

Insert by Queries

Inserting data from queries eliminates the `VALUES` clause—doing so actually replaces the clause with the query. The default and override signatures remain the same. Matching the columns returned by the `SELECT` statement replaces matching `VALUES` clause columns against the override column list or table definition.

The prototype changes slightly for the `INSERT` statement when the source becomes the results from a query rather than a list of values. The `VALUES` keyword must be dropped when you use a query or the `INSERT` statement fails.

The following prototype works for Oracle and MySQL databases and includes options for joining more than one table on various conditions and uses ANSI SQL-92 syntax:

```
INSERT INTO table_name
[( column_name1, column_name2, ...)]
( SELECT    column_value1, column_value2, ...
  FROM      table1
  [[ JOIN table2 ON table1.column_name1 = table2.column_name2] ... ]
  [ WHERE    some_logical_conditions ] );
```


The prototype provides an optional list of column names for an override signature. Then it provides a single table query, optional join with potentially more joins, and ultimately a `WHERE` clause to filter rows. The column names of the query do not need to match the column names of the table to which the values are being added. Only the data types of the query's columns must match the data types of the table's columns. Fixed- and variable-length strings in queries must contain no more than the number of bytes or characters allowed by the table's definition.

`INSERT` statements from queries generally collect information from a series of tables. The number of rows returned for all columns must match. A problem occurs when literal values must map to a multiple row data set. The solution to this problem is to multiply the single row of the literal values against the multiple rows of the data set using a `CROSS JOIN`, which produces a Cartesian Product (every row in one table matched to every row in the other).

For example, let's say you need to enter an active price flag of *Y* with an amount of \$5 and an inactive price flag of *N* with an amount of \$3 for all merchandise in your item table. You need to multiply the single row of numeric and string literals by the number of rows in the table before inserting them into the price table. The solution would be an `INSERT` statement from a query that fabricates (or produces) a data set not found in your raw data, like this:

```
SQL> INSERT INTO price
2  (SELECT    price_s1.NEXTVAL
3    ,
4    i.item_id
5    ,
6    dt.active_flag
7    ,
8    dt.amount
9  FROM      (SELECT    'Y' AS active_flag
10             ,
11             5 AS amount
12            FROM      dual
13            UNION ALL
14            SELECT    'N' AS active_flag
15             ,
16             3 AS amount
17            FROM      dual) dt CROSS JOIN item i);
```

The in-line view creates a two-column by two-row table with the values 'Y' and 5 in one row and the values 'N' and 3 in the other row. The number of rows in the item table will be multiplied by the two rows in the fabricated table and the combined and balanced row set joined by a sequence value for each row in the outer `SELECT` clause. The only change required to implement this sample in a MySQL database switches the `price_s1.NEXTVAL` pseudo column with a null value.

**NOTE**

The CROSS JOIN is a forward reference to Chapter 11. If you're new to SQL and want to understand the full nature of the Oracle and MySQL examples, you should check out the basics of join mechanics in Chapter 11.

The next two sections demonstrate the syntax for INSERT statements with queries, with two approaches shown: the first shows a query that contains all the mandatory columns, and the other shows a query that contains only some of the mandatory columns with additional values that come from string literals. The only differences between Oracle and MySQL are the functions that can convert a query return column to match the definition of a table column and the handling of automatic numbering through sequences.

Oracle Insert by Queries

Here we'll show a query that contains all the columns necessary to insert a row in a table and a query that doesn't contain all the columns necessary for the INSERT statement. The latter mixes string literals with query results through a CROSS JOIN operation, which fabricates data sets.

The INSERT statement will put the data into a table defined as follows:

Name	Null?	Type
-----	-----	-----
KNIGHT_ID	NOT NULL	NUMBER
KNIGHT_NAME	NOT NULL	VARCHAR2 (20)
KINGDOM_ALLEGIANCE_ID	NOT NULL	NUMBER

The knight table has three columns. The knight_id column is first and holds a surrogate key column that would store a sequence value. The knight_name and kingdom_allegiance_id columns hold a natural key, which is unique. That means there can be only one *Peter the Magnificent* with allegiance to a given kingdom stored in the kingdom table, because the kingdom_allegiance_id column holds a foreign key that points to the kingdom_id column in the kingdom table. The following conventions are used in this model: Primary key columns are surrogate key columns, and they use the name of the table plus an _id suffix; foreign key columns can have names that differ from their related primary key column names.

The query finds knights in the available_knight table. The following INSERT statement puts the data into the table:

```
SQL> INSERT INTO knight
2  ( knight_id
3  , knight_name
4  , kingdom_allegiance_id )
5  ( SELECT  knight_s1.NEXTVAL
6  ,
7  ,          (SELECT  kingdom_id
8  ,              FROM    kingdom
9  ,              WHERE   kingdom_name = 'Narnia')
10 FROM      available_knight);
```

The query gets the values from the `available_knight` table, and a subquery on lines 5 through 10 gets the correct foreign key value from the `kingdom` table. Inside the query, a call to the `knight_s1` sequence value ensures the surrogate key is automatically populated.

An alternative syntax would exclude the surrogate key column, because an `ON INSERT` database trigger populates the column. Database triggers support this activity effectively, and many developers use them for this purpose.

The use of an override signature means you can reshuffle the order of columns to fit particular business needs. Be sure to make the same changes in the override signature that you make in the query. The number of possibilities varies with the complexity of queries required to solve the business problems.

MySQL Insert by Queries

Leveraging the same example from the “Oracle Insert by Queries” section, you need only change the first column value returned by the query. The `.NEXTVAL` pseudo column is an Oracle-only solution. In MySQL, you would perform the query with a `null` value as the first column of the query, like so:

```
mysql> INSERT INTO knight
-> ( knight_id
-> , knight_name
-> , kingdom_allegiance_id )
-> ( SELECT null
-> , knight_name
-> , (SELECT kingdom_id
-> FROM kingdom
-> WHERE kingdom_name = 'Narnia')
-> FROM available_knight);
```

An alternative example would implement an override signature that excludes the surrogate key (or auto-incrementing) column. The syntax would then be as follows:

```
mysql> INSERT INTO knight
-> ( knight_name
-> , kingdom_allegiance_id )
-> ( SELECT knight_name
-> , (SELECT kingdom_id
-> FROM kingdom
-> WHERE kingdom_name = 'Narnia')
-> FROM available_knight);
```

The power of inserting from queries is substantial. It lets you collect data from many sources and then insert rows of data in a single statement.

Summary

In this chapter, you learned about inserting data. You can insert a row at a time with a `VALUES` clause and many rows at a time with a subquery. The syntax for these statements differs little between the Oracle and MySQL databases. The most significant issue is handling sequence values.

Mastery Check

The mastery check is a series of true or false and multiple choice questions that let you confirm how well you understand the material in the chapter. You may check the Appendix for answers to these questions.

1. **True** ☐ **False** ☐ An INSERT statement supports multiple row inserts through the VALUES clause in Oracle.
2. **True** ☐ **False** ☐ An INSERT statement supports multiple row inserts through the VALUES clause in MySQL.
3. **True** ☐ **False** ☐ The list of columns in an override signature must match the number and data types of the list of values in the VALUES clause.
4. **True** ☐ **False** ☐ You can insert data into nested tables with an INSERT statement in an Oracle database.
5. **True** ☐ **False** ☐ In an Oracle database, you can use a .CURRVAL pseudo column in a VALUES clause provided the sequence has been placed in scope through an earlier .NEXTVAL pseudo column call.
6. **True** ☐ **False** ☐ A null value that maps to an auto-incrementing column in a MySQL database inserts the next number in the table's sequence.
7. **True** ☐ **False** ☐ In MySQL, you can use an override signature only when you want to insert sequence values.
8. **True** ☐ **False** ☐ An INSERT statement that uses a query can insert one to many rows of data in both Oracle and MySQL databases.
9. **True** ☐ **False** ☐ MySQL supports the .NEXTVAL pseudo column.
10. **True** ☐ **False** ☐ Both Oracle and MySQL databases let you use subqueries in the VALUES clause of an INSERT statement.
11. Which of the following data types requires a built-in function call to put a long variable length character string in a VALUES clause of an INSERT statement?
 - A. BLOB
 - B. MEDIUMCLOB
 - C. TEXT
 - D. CLOB
 - E. VARCHAR
12. In an Oracle database, how many rows can you insert through a VALUES clause?
 - A. 1
 - B. 2
 - C. 3
 - D. 4
 - E. Many

13. In an MySQL database, how many rows can you insert through a `VALUES` clause?
 - A. 1
 - B. 2
 - C. 3
 - D. 4
 - E. Many
14. When you use a query instead of a `VALUES` clause, which of the following is true about the query?
 - A. The query can contain scalar subqueries.
 - B. The query can return only one row at a time.
 - C. The query must return a unique data set.
 - D. The query can't be the product of a join between two or more tables.
 - E. The query is independent of the table structure, which means you can return more columns than the table will accept.
15. Which of the following can't be put in a `VALUES` clause?
 - A. 1
 - B. 2
 - C. 3
 - D. Unlimited
 - E. None of the above